

Communication-Efficient Learning of Deep Networks from Decentralized Data: A Literature-Review Entry

Literature review entry

arXiv:1602.05629

Abstract

McMahan et al. [2017] introduce *Federated Learning* and the FEDERATEDAVERAGING (FedAvg) algorithm: a shared model is trained from data that never leaves client devices by having each client run several epochs of local stochastic gradient descent and a coordinating server average the resulting *models*. The paper’s stated contribution is not a new optimizer—iterative model averaging predates it [McDonald et al., 2010, Povey et al., 2015, Zhang et al., 2015]—but the identification and formalization of *federated optimization*, a regime defined by non-IID, unbalanced, massively distributed, and communication-limited data that violates every standing assumption of the convex distributed-optimization literature. Across five architectures and four datasets, FedAvg reduces the communication rounds required to reach a target accuracy by 10–100× relative to synchronized SGD, establishing the empirical viability of on-device collaborative training and seeding the secure-aggregation and differential-privacy lines that followed.

Notation

The full notation reference is the [math-foundations glossary](#). The symbols this paper introduces or specialises are below (the glossary was network-unreachable at study time, so each row is flagged for upstream propagation):

Symbol	Read aloud as	Meaning
$w \in \mathbb{R}^d$	“w in R-d”	Model parameter vector
$f(w)$	“f of w”	Global objective: mean loss over all n examples
$f_i(w) = \ell(x_i, y_i; w)$	“f-sub-i”	Per-example loss on (x_i, y_i)
K	“K”	Total number of clients
$P_k, n_k = P_k $	“P-sub-k”	Index set / example count on client k
$F_k(w)$	“cap-F-sub-k”	Client k ’s local objective (mean loss on P_k)
$g_k = \nabla F_k(w_t)$	“g-sub-k”	Client k ’s local average gradient at w_t
C, E, B	“C, E, B”	Client fraction / local epochs / local minibatch size
$S_t, m_t = \sum_{k \in S_t} n_k$	“S-sub-t”	Selected clients in round t / their total examples
$u_k = En_k/B$	“u-sub-k”	Local SGD updates client k performs per round

1 Background and Motivation

Mobile devices collectively hold large quantities of privacy-sensitive data—typed text, photographs, sensor traces—that could train high-quality models but cannot responsibly be centralized. McMahan et al. [2017] frame the resulting learning problem as *federated optimization* and distinguish

it from data-center distributed optimization by four properties: the data is *non-IID* (a client’s examples reflect one user and need not resemble the population), *unbalanced* (clients hold very different amounts of data), *massively distributed* (clients far outnumber the examples per client), and subject to *limited communication* (devices are intermittently available on slow links). Prior communication-efficient distributed optimization [Zhang et al., 2012, Zinkevich et al., 2010] assumes the negation of these: many examples per node, IID partitions, and equal shard sizes. The economic premise of the paper is an inversion of the data-center cost model: communication, not computation, dominates, so the goal becomes to *spend additional on-device computation to reduce the number of communication rounds*. Earlier iterative-averaging methods [McDonald et al., 2010, Povey et al., 2015, Zhang et al., 2015] had been studied only in the cluster setting (≤ 16 workers, IID), and asynchronous approaches [Dean et al., 2012] require prohibitively many updates in the federated regime; FedAvg adapts the averaging idea to the federated setting and supplies the appropriate evaluation methodology.

2 Technical Approach

FedAvg targets any finite-sum objective. With n examples partitioned across K clients, client k holding index set P_k with $n_k = |P_k|$, the global objective decomposes exactly into a data-size-weighted average of per-client objectives.

Definitions. The per-example loss is $f_i(w) = \ell(x_i, y_i; w)$ and the per-client objective is $F_k(w) = \frac{1}{n_k} \sum_{i \in P_k} f_i(w)$. Because $\{P_k\}$ partitions the data, $\sum_k n_k = n$, so the weights n_k/n form a probability vector. The setting is IID when a uniformly random partition gives $\mathbb{E}_{P_k}[F_k(w)] = f(w)$; the non-IID setting is defined by the failure of this identity (F_k may be an arbitrarily bad approximation to f).

Central object. The central quantity studied is the finite-sum objective written in client-decomposed form,

$$\min_{w \in \mathbb{R}^d} f(w), \quad f(w) = \frac{1}{n} \sum_{i=1}^n f_i(w) = \sum_{k=1}^K \frac{n_k}{n} F_k(w), \quad (1)$$

where the second equality is an exact algebraic identity, valid for *any* partition (it never invokes the IID assumption).

Method. Differentiating (1) gives the gradient-aggregation identity $\sum_k \frac{n_k}{n} g_k = \nabla f(w_t)$ with $g_k = \nabla F_k(w_t)$. The baseline FEDSGD ($B = \infty$, $E = 1$) therefore performs, with full client participation ($C = 1$), exact distributed full-batch gradient descent,

$$w_{t+1} \leftarrow w_t - \eta \sum_{k=1}^K \frac{n_k}{n} g_k = w_t - \eta \nabla f(w_t). \quad (2)$$

A one-line algebraic fact—that averaging a single local gradient step over clients equals one global gradient step, because $\sum_k n_k/n = 1$ —lets the computation be reorganized into *model averaging*: clients step locally, the server averages models. FedAvg generalizes this by iterating the local update for $u_k = En_k/B$ steps before averaging; the control parameters are C (client fraction), E (local epochs), and B (local minibatch size). Once more than one local step is taken the model-average is no longer a gradient step on f , and that gap—to leading order the client-weighted covariance

$\eta^2 \text{Cov}_k(H_k, g_k)$ between local Hessians and gradients, a pure *heterogeneity* term—is exactly the extra progress FedAvg buys per round. The server aggregation, corrected by an erratum in the published version, averages over the selected set S_t with the participating mass $m_t = \sum_{k \in S_t} n_k$,

$$w_{t+1} \leftarrow \sum_{k \in S_t} \frac{n_k}{m_t} w_{t+1}^k, \quad (3)$$

which is a genuine convex combination for every realization; normalizing by n instead of m_t (the original error) shrinks the aggregate by the participation fraction C each round. Because parameter-space averaging of non-convex networks can in general produce an arbitrarily bad model, the method relies on a shared per-round initialization: McMahan et al. [2017] show empirically (their Figure 1) that averaging two networks trained from *independent* initializations crosses a loss barrier, whereas averaging from a *shared* initialization yields a model better than either parent. Broadcasting the same w_t to all clients each round engineers the favorable regime.

3 Key Results

The evaluation spans the MNIST 2NN (199,210 parameters) and CNN (1,663,370), a character-level Shakespeare LSTM (866,578), a CIFAR-10 CNN ($\sim 10^6$), and a 4,950,544-parameter word-level LSTM trained on over 500,000 clients. Holding the client fraction at $C = 0.1$ and increasing the per-round local computation (larger E , smaller B , hence larger $u = nE/(KB)$) reduces the rounds to a target accuracy by up to $35\times$ for the MNIST CNN and $46\times$ for the 2NN under IID partitions. Under the deliberately *pathological* non-IID partition—data sorted by label into 200 shards of 300, two shards per client, so most clients see only two digits—the speedups shrink to $2.8\text{--}3.7\times$ but remain positive, which the authors read as evidence of robustness. The largest reported gain, $95\times$, occurs on the *unbalanced* non-IID Shakespeare task, which the authors note is actually *easier* because some roles carry large local datasets. On CIFAR-10, FedAvg reaches 85% accuracy in 2,000 rounds versus centralized SGD’s 197,500 minibatch updates, and on the large-scale next-word task FedAvg reaches the target in 35 rounds versus FedSGD’s 820 ($23\times$). The authors further conjecture a regularization benefit akin to dropout [Srivastava et al., 2014], and observe that for very large E the method can plateau or diverge, motivating a learning-rate-style decay of local computation.

4 Position in the Literature

The algorithmic lineage of FedAvg is explicitly prior art: iterative model averaging was studied for the structured perceptron by McDonald et al. [2010], for speech DNNs by Povey et al. [2015], and asynchronously by Zhang et al. [2015], all in the cluster regime. The aggressive endpoint of the FedAvg family—one-shot averaging—is known to be worst-case no better than single-client training in the convex IID case [Zhang et al., 2012, Zinkevich et al., 2010], which is precisely why iteration and a shared initialization matter. The justification that shared-initialization averaging behaves well leans on the over-parameterized loss-landscape results of Goodfellow et al. [2015], Dauphin et al. [2014], and Choromanska et al. [2015]; the underlying obstruction it sidesteps—permutation symmetry rendering independently trained networks incompatible under averaging—is the same phenomenon later formalized by Ainsworth et al. [2023]. The synchronous large-batch baseline (FedSGD) is taken from Chen et al. [2016]. On the privacy axis, the paper offers a qualitative data-minimization argument rather than a guarantee, and forward-references the secure-aggregation protocol of Bonawitz et al. [2017] and the differential-privacy machinery of Abadi et al. [2016]; the communication-reduction techniques of Konečný et al. [2016] are its immediate companion. The

heterogeneity gap that (3) and the covariance term expose is the object that subsequent federated-optimization theory was built to control—via a proximal term in FedProx [Li et al., 2020] and via control variates in SCAFFOLD [Karimireddy et al., 2020].

5 Limitations and Open Questions

The paper proves nothing about FedAvg: there is no convergence theorem, no rate, and no condition guaranteeing non-divergence—the authors concede the method “can plateau or diverge” for large E , and that parameter averaging “could produce an arbitrarily bad model.” The equivalence to gradient descent holds only at the FedSGD endpoint ($E = 1$); for $E > 1$ no fixed point is characterized. The evaluation is structurally favorable to the proposed method: the learning rate is grid-searched per x -axis point over 11–13 values (more than 2,000 models trained), learning curves are monotonized, round counts are obtained by linear interpolation, and no variance or confidence intervals are reported. The FedSGD baseline—one full-batch gradient step per round—is the weakest possible per-round comparator, so a large per-round speedup is partly definitional. The non-IID robustness claim rests on a single artificial MNIST partition plus a proprietary, unverifiable social-media corpus, and several non-IID configurations are in fact *slower* than FedSGD. The load-bearing premise that on-device “computation is essentially free” relative to communication is asserted rather than measured, despite up to $u_k = 1200$ local updates per round. Finally, the published erratum—normalizing the aggregate by m_t rather than n in (3)—means the corrected estimator is only *exactly* unbiased when clients are balanced; for unbalanced n_k under partial participation it is a ratio estimator with $O(1/m)$ bias, a subtlety the paper does not analyze. The central open questions are therefore: what fixed point does FedAvg approach under non-IID data, is there a principled rule for choosing (E, B) , and what does it take to convert the qualitative privacy story into a formal guarantee.

Sandbox Probe

A companion sandbox reproduces the two load-bearing claims at toy scale. A pure-NUMPY federated softmax-regression demo (`toy_fedavg.py`) recovers the qualitative shape of the paper’s Table 2: with $K = 100$, $C = 0.1$, the best FedAvg configuration ($E = 20$, $B = 10$) reaches the target accuracy in 2 communication rounds versus FedSGD’s 62 ($31\times$), with strictly smaller speedups under the pathological non-IID partition. A second numpy demo (`tiny_fedavg_averaging.py`) reproduces the Figure 1 phenomenon on a one-hidden-layer MLP: averaging from independent initializations yields a positive loss barrier (+0.024), while shared-initialization averaging yields a model below both parents (−0.027). A PyTorch script reproduces the result on the paper’s exact MNIST CNN (1,663,370 parameters) on Apple-Silicon MPS, and a PEFT-based script applies FedAvg to LoRA adapters, communicating only 4.3% of parameters per round—a forward-looking instance of the same algorithm.

References

Martin Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2016.

- Samuel K. Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. Git re-basin: Merging models modulo permutation symmetries. In *International Conference on Learning Representations (ICLR)*, 2023.
- Keith Bonawitz, Vladimir Ivanov, Ben Kreuter, Antonio Marcedone, H. Brendan McMahan, Sarvar Patel, Daniel Ramage, Aaron Segal, and Karn Seth. Practical secure aggregation for privacy-preserving machine learning. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, 2017.
- Jianmin Chen, Rajat Monga, Samy Bengio, and Rafal Jozefowicz. Revisiting distributed synchronous sgd. In *ICLR Workshop Track*, 2016.
- Anna Choromanska, Mikael Henaff, Michaël Mathieu, Gérard Ben Arous, and Yann LeCun. The loss surfaces of multilayer networks. In *Proceedings of the 18th International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2015.
- Yann N. Dauphin, Razvan Pascanu, Caglar Gulcehre, Kyunghyun Cho, Surya Ganguli, and Yoshua Bengio. Identifying and attacking the saddle point problem in high-dimensional non-convex optimization. In *Advances in Neural Information Processing Systems (NIPS)* 27, 2014.
- Jeffrey Dean, Greg S. Corrado, Rajat Monga, Kai Chen, Matthieu Devin, Quoc V. Le, Mark Z. Mao, Marc’Aurelio Ranzato, Andrew Senior, Paul Tucker, Ke Yang, and Andrew Y. Ng. Large scale distributed deep networks. In *Advances in Neural Information Processing Systems (NIPS)* 25, 2012.
- Ian J. Goodfellow, Oriol Vinyals, and Andrew M. Saxe. Qualitatively characterizing neural network optimization problems. In *International Conference on Learning Representations (ICLR)*, 2015.
- Sai Praneeth Karimireddy, Satyen Kale, Mehryar Mohri, Sashank J. Reddi, Sebastian U. Stich, and Ananda Theertha Suresh. SCAFFOLD: Stochastic controlled averaging for federated learning. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020.
- Jakub Konečný, H. Brendan McMahan, Felix X. Yu, Peter Richtárik, Ananda Theertha Suresh, and Dave Bacon. Federated learning: Strategies for improving communication efficiency. In *NIPS Workshop on Private Multi-Party Machine Learning*, 2016.
- Tian Li, Anit Kumar Sahu, Manzil Zaheer, Maziar Sanjabi, Ameet Talwalkar, and Virginia Smith. Federated optimization in heterogeneous networks. In *Proceedings of Machine Learning and Systems (MLSys)*, 2020.
- Ryan McDonald, Keith Hall, and Gideon Mann. Distributed training strategies for the structured perceptron. In *Human Language Technologies: The 2010 Annual Conference of the North American Chapter of the ACL (NAACL HLT)*, 2010.
- H. Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Agüera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, volume 54 of *PMLR*, 2017. URL <https://arxiv.org/abs/1602.05629>.
- Daniel Povey, Xiaohui Zhang, and Sanjeev Khudanpur. Parallel training of deep neural networks with natural gradient and parameter averaging. In *ICLR Workshop Track*, 2015.

- Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: A simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15, 2014.
- Sixin Zhang, Anna Choromanska, and Yann LeCun. Deep learning with elastic averaging sgd. In *Advances in Neural Information Processing Systems (NIPS)* 28, 2015.
- Yuchen Zhang, Martin J. Wainwright, and John C. Duchi. Communication-efficient algorithms for statistical optimization. In *Advances in Neural Information Processing Systems (NIPS)* 25, 2012.
- Martin Zinkevich, Markus Weimer, Lihong Li, and Alex J. Smola. Parallelized stochastic gradient descent. In *Advances in Neural Information Processing Systems (NIPS)* 23, 2010.